

1)Feature extraction from entire dataset

extract_environmental_features_without_labels.py (uses helper script-tc_binary_classification_helpers)

Extracting features for whole world,data saved in **data/extracted_features_entire_world**

PYTHONPATH="Not our/scripts" python3 extract_environmental_features_without_labels.py "data/entire dataset in grib" data/extracted_features_entire_world --processes 14

2)inspect using check_features.py

===== GEOGRAPHIC COVERAGE =====

Latitude Bounds : -90.0° to 90.0°

Latitude Steps : 181 points

Longitude Bounds : 0.0° to 359.0°

Longitude Steps : 360 points

===== EXTRACTED VARIABLES =====

['ugrdprs', 'vgrdprs', 'vvelprs', 'tmpprs', 'rhprs', 'hgtprs', 'absvprs', 'pressfc', 'capesfc', 'tmprfc', 'landmask']

===== VERTICAL LEVELS =====

[np.float64(1000.0), np.float64(975.0), np.float64(950.0), np.float64(925.0), np.float64(900.0), np.float64(850.0), np.float64(800.0), np.float64(750.0), np.float64(700.0), np.float64(650.0), np.float64(600.0), np.float64(550.0), np.float64(500.0), np.float64(450.0), np.float64(400.0), np.float64(350.0), np.float64(300.0), np.float64(250.0), np.float64(200.0)]

(base) puranjay@PC-RRSCN-05:/mnt/e/tc-dl\$

3)Label with ibtracs(create_labels.py)

Key Corrections & Scientific Rigor Applied:

1. Full Active TC Tracking: Captures all active TCs at each timestamp (**Active TC Locations** and **Active TC SIDs**), regardless of whether a new genesis event is occurring.
2. Proper ISO Formatting: Ensures **Other TC Locations** and **Active TC Locations** are cleanly serialized as standard JSON strings/lists (e.g., `[[lat, lon], ...]`) so **pandas** can read them back reliably without string-mangling.

3. Multi-Genesis Support: If multiple cyclogenesis events occur simultaneously at the same timestamp across the globe, it records all genesis coordinates as a list rather than overwriting or taking only the first one.
4. Clean Domain & Coordinate Mapping: Keeps 0°–360° global longitude alignment and accurately records lead times without dropping non-genesis atmospheric states.

```

--> Listing reanalysis NetCDF files from: data/extracted_features_entire_world
--> Extracted 12827 NetCDF observation timestamps.
--> Dataset Domain Lat Bounds: [-90.0°, 90.0°], Lon Bounds: [0.0°, 359.0°]
--> Loading IBTrACS best-track dataset from: data/ibtracs/ibtracs.ALL.list.v04r01.csv
Generating Labels: 100%|██████████|

Label generation complete!
  • Total Observation Timestamps : 12827
  • Positive Genesis Timestamps : 1026
  • Output Saved To : data/extracted_features_entire_world/tc_48h.csv
PS E:\TC-DL> docker run --rm -it --gpus all `
>> -v /usr/lib/wsl:/usr/lib/wsl `
>> -e NVIDIA_VISIBLE_DEVICES=all `
>> -e NVIDIA_DRIVER_CAPABILITIES=all `
>> -e LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" `
>> -e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
>> -e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
>> -v "${PWD}:/workspace" `
>> tc-dl-pipeline:latest /workspace/scripts/create_labels1.py `
>> --best-track "data/ibtracs/ibtracs.ALL.list.v04r01.csv" `
>> --observations-dir "data/extracted_features_entire_world" `
>> --leadtime 72
--> Listing reanalysis NetCDF files from: data/extracted_features_entire_world
--> Extracted 12827 NetCDF observation timestamps.
--> Dataset Domain Lat Bounds: [-90.0°, 90.0°], Lon Bounds: [0.0°, 359.0°]
--> Loading IBTrACS best-track dataset from: data/ibtracs/ibtracs.ALL.list.v04r01.csv
Generating Labels: 100%|██████████|

Label generation complete!
  • Total Observation Timestamps : 12827
  • Positive Genesis Timestamps : 1030
  • Output Saved To : data/extracted_features_entire_world/tc_72h.csv
PS E:\TC-DL>

```

5.

To run: docker run --rm -it --gpus all `

>> -v /usr/lib/wsl:/usr/lib/wsl `

>> -e NVIDIA_VISIBLE_DEVICES=all `

>> -e NVIDIA_DRIVER_CAPABILITIES=all `

>> -e

LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" `

>> -e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `

>> -e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `

>> -v "\${PWD}:/workspace" `

>> tc-dl-pipeline:latest /workspace/scripts/create_labels1.py `

```
>> --best-track "data/ibtracs/ibtracs.ALL.list.v04r01.csv" `
>> --observations-dir "data/extracted_features_entire_world" `
>> --leadtime 72
```

4)Chronological, year-based split using split_data.py

This script divides your data into Train, Validation, and Test datasets based on the timeline we established. It then saves them as three separate lightweight CSV files inside your features directory so they can be fed cleanly into our TensorFlow generator.

Key Improvements & Scientific Rigor Applied

1. **Flexible Column Support:** Supports both `Observation_Date` (from our updated `create_labels.py`) and legacy `Date` columns without failing.
2. **Flexible Lead-Time CLI Support:** Accepts `--csv-path` and `--leadtime` flags so you can split `tc_24h.csv`, `tc_48h.csv`, and `tc_72h.csv` dynamically.
3. **Distribution Audit:** Prints positive vs. negative counts for Train, Validation, and Test splits so you can verify class distributions across years.

```
Splits:
- Train      : 2008 - 2018
- Validation : 2019 - 2020
- Test       : 2021 - 2022
.....
```

Lead Time	Train Positives (2008–2018)	Validation Positives (2019–2020)	Test Positives (2022)
24h	729	155	136
48h	732	158	136
72h	736	159	135

To run:

```
docker run --rm -it --gpus all `
-v /usr/lib/wsl:/usr/lib/wsl `
-e NVIDIA_VISIBLE_DEVICES=all `
-e NVIDIA_DRIVER_CAPABILITIES=all `
-e
LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" `
-e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
-e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
-v "${PWD}:/workspace" `
tc-dl-pipeline:latest /workspace/scripts/split_data.py `
--csv-path "data/extracted_features_entire_world/tc_24h.csv" `
--output-dir "data/extracted_features_entire_world"
```

Split 48h CSV

```
docker run --rm -it --gpus all -v "${PWD}:/workspace" tc-dl-pipeline:latest
/workspace/scripts/split_data.py --csv-path
"data/extracted_features_entire_world/tc_48h.csv" --output-dir
"data/extracted_features_entire_world"
```

Split 72h CSV

```
docker run --rm -it --gpus all -v "${PWD}:/workspace" tc-dl-pipeline:latest
/workspace/scripts/split_data.py --csv-path
"data/extracted_features_entire_world/tc_72h.csv" --output-dir
"data/extracted_features_entire_world"
```

5)Precompute Fixed Global Normalization Statistics

Run a one-time precomputation script to generate `channel_means.npy` and `channel_stdns.npy` over the entire training set (or a uniform sampling across all training NetCDF files). This will be feeded to `tcdl_generator.py`

pipeline.

Computing statistics over polar ice sheets ($> 35^{\circ}\text{N}/\text{S}$) severely skews temperature, geopotential height, and relative humidity distributions away from the tropical atmosphere ($35^{\circ}\text{S} - 35^{\circ}\text{N}$). Furthermore, setting NaNs to `0.0` before computing sums heavily corrupts empirical means (especially for sparse fields like CAPE or relative humidity).

Here is the updated, scientifically rigorous `precompute_stats.py` script incorporating tropical domain filtering, per-channel valid pixel counting, NaN masking, and NetCDF attribute cleaning.

To run: `docker run --rm -it --gpus all ``

`-v /usr/lib/wsl:/usr/lib/wsl ``

`-e NVIDIA_VISIBLE_DEVICES=all ``

`-e NVIDIA_DRIVER_CAPABILITIES=all ``

`-e`

`LD_LIBRARY_PATH="/usr/lib/wsl/lib/./usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" ``

`-e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" ``

`-e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" ``

`-v "${PWD}:/workspace" ``

`tc-dl-pipeline:latest /workspace/scripts/precompute_stats.py --processes 14`

6)Data processing for model loading,data augmentation etc using `tcdl_generator.py`

(`tcdl_generator.py`) —► Reads the split CSVs to dynamically sample patches during training.It contains the `TropicalCyclogenesisGenerator` class definition.

Core Improvements & Scientific Rigor Applied:

1. **Elimination of Fixed (0°, 180°) Negatives:** Negatives are dynamically sampled across the global tropical ocean grid (bounded by 35°S — 35°N and checked against the `Landmask` variable).
2. **Dynamic Exclusion Masking:** Uses the **Haversine formula** to calculate precise distances on the spherical Earth. Any sampled coordinate within 400 km of:
 - Future genesis locations (`Genesis_Lats` , `Genesis_Lons`), **or**
 - Active, existing tropical cyclone tracks (`Active_TC_Locations`),
is rejected on-the-fly to prevent corrupting negative distributions with real storms.
3. **Multi-Negative Ratio:** Dynamically samples a ratio of N global negatives per positive sample per epoch (default ratio = 5 : 1 or 10 : 1, configurable) to reflect realistic oceanic spatial distributions.
4. **Meteorological Augmentations (Domain-Preserving):**
 - **Spatial Jittering:** Randomly shifts patch centers by ± 1 or ± 2 grid cells (70% chance).
 - **Gaussian Field Noise:** Low-magnitude Gaussian noise ($\sigma = 0.005$) on atmospheric fields (50% chance).
 - **Channel Perturbation:** Low-level scaling ($\pm 1\%$) across physical channels (30% chance).
 - **Pixel Masking:** Random 1% spatial pixel masking (20% chance).
 - **Forbidden:** No horizontal/vertical flips or non-90° rotations (preserves Coriolis force physics).
5. **Channel Consistency:** Strictly enforces the 137-channel array structure (7 3D variables $\times$ 19 vertical pressure levels + 4 surface variables).

🔧 Key Fixes Applied in the Redesigned Generator

1. **Fixed Spatial Jitter (Extraction-Based Jitter):**
 - Removed `np.roll(...)` completely to prevent wrapping boundary artifacts.
 - Positives and negatives now perform spatial jittering by dynamically shifting the target extraction center (± 1 or ± 2 grid cells in Lat/Lon) directly during patch extraction from the NetCDF array.
2. **Re-extracted & Dynamic Positives:**
 - Positives are no longer statically cached forever. During `resample_dataset()` , positives are re-extracted with randomized coordinate offsets per epoch so the model never sees fixed, frozen positive crops.
3. **Complete Storm Exclusion (Genesis + Active TCs):**
 - Exclusion masking explicitly combines **Active TCs** (existing mature systems) AND **Target Genesis Locations**. Any candidate point within 400 km of *any* active storm or genesis point is rejected from the negative distribution.
4. **Meteorological-Only Augmentation (Landmask Protected):**
 - Landmask (channel index 136) is kept strictly untouched during augmentation.
 - Noise, channel perturbations, and pixel masking are applied **only** to the first 136 physical meteorological fields (wind, temperature, humidity, surface pressure, CAPE, etc.).
5. **Channel-Relative Noise Scaling:**
 - Gaussian noise is now dynamically scaled relative to the empirical standard deviation (σ_c) of each specific channel ($0.01 \times \sigma_c$) rather than applying a flat 0.005 across vastly different units (Pa, K, m/s, s⁻¹).
6. **Strict Validation Safeguards:**
 - Guarantees `enable_augmentation=False` for validation/test generators.

To run the updated `train_cnn_lstm.py` script cleanly, you need to make **two targeted updates** to `tcdl_generator.py` :

1. **Compute and Apply Per-Channel Standard Normalization (μ_c, σ_c):** Ensure feature variables with drastically different units (Pa, K, m/s, s⁻¹) are normalized on-the-fly before entering the model layers.
2. **Handle NaN / Inf Values Cleanly:** Ensure empty or corrupted grid slices default to normalized zeros rather than unscaled arrays.

Step 2: Complete, Production-Ready `tcdl_generator.py`

This version incorporates:

1. **Fixed Offline Global Stats Loading (`channel_means.npy` & `channel_stds.npy`):** Eliminates single-epoch sampling variance.
2. **Expanded Exclusion Radius (550 km):** Prevents sampling pre-genesis wave perturbations or outer cyclone feeder bands as negatives.
3. **Hard-Negative Bias (70% Active Wave Disturbance / 30% Random Ocean):** Samples negatives preferentially from active convergence zones and tropical wave regions.
4. **Normalized-Space Augmentation:** Performs z-score normalization **first**, applying relative noise ($\sigma = 0.01$) uniformly in standard normalized feature space.
5. **Fixed, Deterministic Validation/Testing:** Guarantees validation sets evaluation remains identical across epochs with zero dynamic sampling jitter or augmentations.

Augmentation:

Here is why there are 729 positive labels in the training set.

1. What 729 Represents (Timesteps, Not Augmentations)

The number **729** is the count of **unique historical tropical cyclogenesis events (genesis timesteps)** in the IBTrACS database across the 11-year training period (2008–2018) for a 24-hour lead time target.

On average, Earth produces ~80–85 tropical cyclones per year globally across all ocean basins:

- $11 \text{ years} \times \approx 65\text{--}70 \text{ genesis events/year} \approx \mathbf{729}$ unique genesis timesteps.

2. How Data Augmentation Works With the Generator

In standard image classification (like training on static JPEG files), augmentation increases the *total dataset size on disk* (e.g., duplicating 729 images into 3,645 static files).

In our **Dynamic Generator** (`tcdl_generator.py`), augmentation happens **on-the-fly in RAM every single epoch**:

- **Epoch 1:** Reads the 729 genesis timesteps → applies spatial jitter (e.g., $-1 \text{ lat}, +2 \text{ lon}$) → adds normalized Gaussian noise → yields **729 unique augmented patches** + 2,187 dynamic ocean negatives (3 : 1 ratio).
- **Epoch 2:** Reads the *same* 729 genesis timesteps → applies *new* spatial jitter (e.g., $+2 \text{ lat}, -1 \text{ lon}$) → adds *new* noise → yields **729 brand new augmented positive crops** + 2,187 new ocean negatives.
- **Epoch N :** The model sees a fresh, newly jittered/perturbed variation of every positive storm center. Over 100 epochs, the model sees $729 \times 100 = 72,900$ **unique spatial variations** of tropical cyclogenesis events!

Applied Augmentations

1. Extraction-Based Spatial Jittering (70% Probability):

- Instead of using `np.roll` (which wraps boundaries across the image edges), the patch center is dynamically shifted by ± 1 or ± 2 grid cells directly during NetCDF array extraction.

2. Normalized-Space Field Noise (50% Probability):

- Low-magnitude Gaussian noise ($\sigma = 0.01$) is injected into the meteorological fields *after* per-channel z-score normalization.

3. Channel Scale Perturbation (30% Probability):

- Multiplicative uniform scaling ($\pm 1\%$) is applied across physical channels to simulate slight observational sensor noise.

4. Random Spatial Pixel Dropout (20% Probability):

- 1% of spatial pixels are randomly masked to encourage feature redundancy.

Crucial Safeguards Included

- **Binary Landmask Preservation:** Augmentations (noise, scaling, and masking) are strictly restricted to the **136 meteorological channels** (channels `0..135`). The binary `landmask` (channel `136`) is kept 100% untouched.
- **Preservation of Atmospheric Physics:** Forbidden augmentations (like non-90° rotations or horizontal/vertical flips) were deliberately excluded to prevent corrupting Coriolis forces and synoptic circulation orientation.
- **Evaluation Safeguard:** Augmentation runs **only during training** (`is_train=True`). The validation and test generators run with `enable_augmentation=False` on fixed evaluation sets.

Key Performance & Architecture Features

1. **Multi-Threaded Parallel Patch Extraction (`ThreadPoolExecutor`):**
 - Uses multi-threaded CPU workers to open NetCDF files and extract patches simultaneously, accelerating dataset generation by $3 \times - 4 \times$.
2. **In-Memory LRU Dataset Handle Cache:**
 - Caches active xarray handles in RAM (`self._dataset_cache`) so files read multiple times during negative sampling do not pay the OS open/close I/O penalty.
3. **Preserved Scientific & Augmentation Rigor:**
 - Hard-negative vorticity mining (70% preference) preserved.
 - Extraction-level spatial jittering ($\pm 1, \pm 2$ grid cells) preserved.
 - Haversine exclusion radius (550 km) preserved.
 - Offline tropical z-score normalization (μ, σ) preserved.
 - Untouched binary landmask channel preserved.

Key Optimizations Added

1. **Thread-Safe LRU Cache (`threading.Lock`):**
 - Multi-threaded workers querying or evicting file handles from `self._dataset_cache` are now protected by a thread-safe mutex lock (`self.cache_lock`), preventing race conditions.
2. **Parallelized Hard-Negative Sampling (`num_workers=16`):**
 - Instead of a single-threaded sequential loop, negative candidate searching and Haversine exclusion checks are distributed across **16 parallel worker threads**.
3. **Pre-Normalized Positive Caching:**
 - Positives are extracted and normalized once in memory during initialization, so each epoch only applies spatial jittering and normalized-space noise—saving unnecessary per-epoch array re-normalizations.
4. **Expanded Cache Size (256 Handles):**
 - Increased handle capacity to 256 active NetCDF files, keeping frequently accessed global daily files directly in RAM.

7)Train cnn_lstm.py

2. **Dynamic Lead-Time Support (`--leadtime 24|48|72`):**
 - The script now dynamically accepts `--leadtime` via CLI and automatically loads the corresponding `tc_{leadtime}h_train.csv` and `tc_{leadtime}h_val.csv` files.
3. **Channel-Wise Standard Normalization (μ, σ):**
 - Integrates `Keras Normalization` across all 137 individual channels (wind, pressure, temperature, humidity, CAPE, etc.) so gradient calculations are well-conditioned across disparate physical units.
4. **Mixed Precision Acceleration (`mixed_float16`):**
 - Enables FP16 tensor core acceleration on your NVIDIA GPU for faster epoch throughput and lower VRAM consumption, while maintaining FP32 precision on the final classification output for numerical stability.
5. **Class Imbalance Loss & PR-AUC Optimization:**
 - Utilizes `BinaryFocalCrossentropy` ($\gamma = 2.0$) or weighted BCE to penalize false negatives heavily.
 - Monitored via **Precision-Recall Area Under the Curve** (`val_pr_auc`) instead of standard `val_loss`.
6. **Gradient Clipping & Reproducibility:**
 - Instantiates Adam with `clipnorm=1.0` to prevent gradient explosions from noisy atmospheric fields.
 - Enforces global random seeding (`random` , `numpy` , `tensorflow`).

How We Make the Spatial CNN-LSTM Physically and Structurally Meaningful

Instead of collapsing everything with `GlobalAveragePooling2D()` :

1. **Spatial Sequence Extraction (Grid/Row Slices as Sequence Steps):**
 - After the 2D convolutions extract high-level feature maps across the patch (e.g., shape: $(5, 5, 128)$ for a 21×21 input), we treat the spatial grid rows ($T = 5$) as a sequential progression across the storm perturbation field.
 - `TimeDistributed` layers process each spatial row/feature vector, converting the spatial patch into a sequence: ($T = 5$, features = 640).
2. **True LSTM Recurrence:**
 - The LSTM layer now processes a real sequence of length $T = 5$ steps across the spatial vorticity/thermodynamic structure, learning directional spatial gradients and circulation structures (like eye formation, feeder bands, or asymmetry).

```

Input: (21, 21, 137)
|
v
Conv2D + BatchNorm + ReLU (64)  → Output: (21, 21, 64)
|
MaxPool2D (2, 2)                → Output: (10, 10, 64)
|
Conv2D + BatchNorm + ReLU (128) → Output: (10, 10, 128)
|
MaxPool2D (2, 2)                → Output: (5, 5, 128)
|
Reshape to Sequence (5, 640)    → (T=5 steps, 5×128=640 spatial feature channels)
|
LSTM (128 units, return_sequences=False) → Processes spatial structure sequentially
|
Dense (128) + Dropout (0.4)
|
Dense (64) + Dropout (0.3)
|
Sigmoid Output (FP32)

```

1. **Dual-Axis Bidirectional Spatial Sequence Modeling:** Processes both row-wise and column-wise spatial feature progressions in parallel before feeding them to the LSTM, eliminating directional bias across the 21×21 cyclone patch.
2. **TensorBoard Logging:** Automatically creates logs in `logs/{leadtime}h/` for tracking PR-AUC, ROC-AUC, Loss, and Learning Rate curves during training.
3. **JSON Config Serialization:** Saves exact hyperparameter configurations to `saved_models/config_cnn_lstm_{leadtime}h.json` alongside the model checkpoint for 100% reproducibility.
4. **Preserved Architecture:** Integrates directly with `tcdl_generator.py`, retaining FP16 Mixed Precision, Binary Focal Crossentropy Loss ($\gamma = 2.0, \alpha = 0.25$), `val_pr_auc` monitoring, gradient clipping (`clipnorm=1.0`), and deterministic random seeds.

```
2026-07-23 05:49:52.088525: I tensorflow/core/common_runtime/gpu/gpu_device.cc:2019] Could not identify NUMA node of platform GPU id 0, defaulting to
2026-07-23 05:49:52.088642: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/de
Your kernel may have been built without NUMA support.
2026-07-23 05:49:52.090504: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1928] Created device /job:localhost/replica:0/task:0/device:GPU:0 with
0:01:00.0, compute capability: 7.5
Model: "CNN_LSTM_Meaningful_Model"
```

Layer (type)	Output Shape	Param #	Connected to
Input_Patch (InputLayer)	(None, 21, 21, 137)	0	-
conv2d (Conv2D)	(None, 21, 21, 64)	78,976	Input_Patch[0][0]
batch_normalization (BatchNormalization)	(None, 21, 21, 64)	256	conv2d[0][0]
activation (Activation)	(None, 21, 21, 64)	0	batch_normalization[0][0]
max_pooling2d (MaxPooling2D)	(None, 10, 10, 64)	0	activation[0][0]
conv2d_1 (Conv2D)	(None, 10, 10, 128)	73,856	max_pooling2d[0][0]
batch_normalization_1 (BatchNormalization)	(None, 10, 10, 128)	512	conv2d_1[0][0]
activation_1 (Activation)	(None, 10, 10, 128)	0	batch_normalization_1[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 128)	0	activation_1[0][0]
Transpose_Columns (Permute)	(None, 5, 5, 128)	0	max_pooling2d_1[0][0]
Row_Spatial_Sequence (Reshape)	(None, 5, 640)	0	max_pooling2d_1[0][0]
Col_Spatial_Sequence (Reshape)	(None, 5, 640)	0	Transpose_Columns[0][0]
Row_LSTM (LSTM)	(None, 64)	180,480	Row_Spatial_Sequence[0][0]
Col_LSTM (LSTM)	(None, 64)	180,480	Col_Spatial_Sequence[0][0]
Dual_Axis_Merge (Concatenate)	(None, 128)	0	Row_LSTM[0][0], Col_LSTM[0][0]
dropout (Dropout)	(None, 128)	0	Dual_Axis_Merge[0][0]
dense (Dense)	(None, 128)	16,512	dropout[0][0]
batch_normalization_2 (BatchNormalization)	(None, 128)	512	dense[0][0]
dropout_1 (Dropout)	(None, 128)	0	batch_normalization_2[0][0]
dense_1 (Dense)	(None, 64)	8,256	dropout_1[0][0]
batch_normalization_3 (BatchNormalization)	(None, 64)	256	dense_1[0][0]
dropout_2 (Dropout)	(None, 64)	0	batch_normalization_3[0][0]
Output_Probability (Dense)	(None, 1)	65	dropout_2[0][0]

Trainable params: 539,393 (2.06 MB)

Non-trainable params: 768 (3.00 KB)

Non-trainable params: 755 (3100 KB)

Extracting Positives (Train): 21%

8)Hybrid Conv-ViT-BiGRU using train_vit_gru.py

a deep, stable Vision Transformer.

Key Architectural & Training Upgrades

1. **Pre-LN Transformer Encoder Stack (4 Blocks):**
 - Expanded from a single shallow attention layer into **4 stacked Transformer Encoder blocks** using modern Pre-LayerNormalization (`LayerNorm` → `MultiHeadAttention` → Residual) for optimal gradient flow.
2. **Enhanced Embedding & Multi-Head Attention:**
 - Increased embedding dimension to `hidden_dim = 192` with `num_heads = 6` (`key_dim = 32`), perfectly matching the multi-channel 137-variable meteorological patch size.
 - Added attention dropout (`dropout=0.1`) inside the `MultiHeadAttention` computation to prevent attention-weight over-concentration.
3. **Optimized AdamW + Focal Loss + FP16 Mixed Precision:**
 - Uses `AdamW` optimizer with `learning_rate=1e-4` and `weight_decay=1e-4`.
 - Compiled with `BinaryFocalCrossentropy(gamma=2.0, alpha=0.25)` to handle class imbalance.
 - Enables FP16 Mixed Precision acceleration for high throughput on GPU tensor cores.
4. **PR-AUC Monitoring & Full Infrastructure Parity:**
 - Monitors `val_pr_auc` for Early Stopping and Model Checkpoints.
 - Supports dynamic CLI lead times (`--leadtime 24|48|72`).
 - Generates TensorBoard logs in `logs/vit_gru_{leadtime}h/` and saves configuration metadata JSON files to `saved_models/config_vit_gru_{leadtime}h.json`.

That is a spot-on critique. It pinpoints the fundamental bottleneck in pure Vision Transformers when applied to 3D meteorological tensor grids: **forcing a linear projection layer to collapse a dense $3 \times 3 \times 137$ local volume (1,233 parameters) into a single 192-dim vector forces the model to learn spatial locality and channel importance from scratch.**

When CNN-LSTMs win on meteorological grids, it is precisely because 2D convolutions provide a built-in spatial inductive bias, naturally extracting local vorticity gradients, thermal cores, and vertical shear structures across all 136 pressure/surface channels.

To give the ViT model every chance of matching or exceeding the CNN-LSTM's `0.9818 ROC-AUC` and `0.9828 PR-AUC`, we need to build a **Hybrid Channel-Attention CNN Stem + Pre-LN Vision Transformer + Bidirectional GRU Architecture**.

Key Upgrades Incorporated

1. Channel-Attention SE Block (Squeeze-and-Excitation)

Before patch extraction, the 137 input channels pass through a lightweight 1D Squeeze-and-Excitation block to learn non-linear channel inter-dependencies. The model explicitly scales physically critical variables (vorticity, vertical velocity, humidity) while dampening background noise.

2. Shallow 2D CNN Feature Stem

Instead of directly slicing raw grid values into linear patches, a 2-layer Conv2D stem (3×3 filters with ReLU and BatchNorm) processes local spatial structures first. This preserves local atmospheric physics (shear lines, vorticity cores) before feeding tokens to the Transformer.

3. $4\times$ Feed-Forward Network (FFN) Expansion Ratio

Expanded the inner Transformer FFN dense layers from `hidden_dim * 2` (384) to standard Transformer capacity: `hidden_dim * 4` (768).

4. Spatial Sequence Bidirectional GRU (BiGRU)

Replaced the unidirectional 128-unit GRU with a `Bidirectional(GRU(128))` layer. This processes patch token sequences in both directions simultaneously across the spatial feature map grid, mirroring the dual-axis spatial coverage of the CNN-LSTM.

5. Standardized Binary Cross-Entropy Loss

Replaced `BinaryFocalCrossentropy` with standard `BinaryCrossentropy()`. Since the dynamic generator already enforces a 3 : 1 hard-negative ocean balance during training, standard cross-entropy provides smoother gradient scaling and prevents under-predicting positive genesis recall.

Your kernel may have been built without NMA support.
2026-07-23 11:02:49.754390: I tensorflow/core/common_runtime/gpu/gpu_device.cc:2019] Could not identify NMA node of platform GPU id 0, defaulting to 0. Your kernel may not
2026-07-23 11:02:49.754390: I tensorflow/core/common_runtime/gpu/gpu_device.cc:2019] Could not identify NMA node of platform GPU id 0, defaulting to 0. Your kernel may not
Your kernel may have been built without NMA support.
2026-07-23 11:02:49.754390: I tensorflow/core/common_runtime/gpu/gpu_device.cc:2019] Could not identify NMA node of platform GPU id 0, defaulting to 0. Your kernel may not
Model: "Hybrid_Conv_ViT_BiGRU_Model"

Layer (type)	Output Shape	Param #	Connected to
Input_Patch (InputLayer)	(None, 21, 21, 137)	0	-
Channel_Attention (SqueezeAndExcitationBlock)	(None, 21, 21, 137)	4,658	Input_Patch[0][0]
Stem_Conv1 (Conv2D)	(None, 21, 21, 64)	78,976	Channel_Attention[0][0]
Stem_BN1 (BatchNormalization)	(None, 21, 21, 64)	256	Stem_Conv1[0][0]
Stem_ReLU1 (Activation)	(None, 21, 21, 64)	0	Stem_BN1[0][0]
Stem_Conv2 (Conv2D)	(None, 21, 21, 128)	73,856	Stem_ReLU1[0][0]
Stem_BN2 (BatchNormalization)	(None, 21, 21, 128)	512	Stem_Conv2[0][0]
Stem_ReLU2 (Activation)	(None, 21, 21, 128)	0	Stem_BN2[0][0]
Patch_Embedder (PatchExtractorAndEmbedder)	(None, 49, 192)	230,784	Stem_ReLU2[0][0]
Transformer_Norm1_Block1 (LayerNormalization)	(None, 49, 192)	384	Patch_Embedder[0][0]
MultiHeadAttention_Block1 (MultiHeadAttention)	(None, 49, 192)	148,224	Transformer_Norm1_Block1[...] Transformer_Norm1_Block1[...]
Attn_Dropout_Block1 (Dropout)	(None, 49, 192)	0	MultiHeadAttention_Block1[...]
Add_Attn_Block1 (Add)	(None, 49, 192)	0	Patch_Embedder[0][0], Attn_Dropout_Block1[0][0]
Transformer_Norm2_Block1 (LayerNormalization)	(None, 49, 192)	384	Add_Attn_Block1[0][0]
FFN_Expansion_Block1 (Dense)	(None, 49, 768)	148,224	Transformer_Norm2_Block1[...]
FFN_Dropout_Block1 (Dropout)	(None, 49, 768)	0	FFN_Expansion_Block1[0][0]
FFN_Projection_Block1 (Dense)	(None, 49, 192)	147,648	FFN_Dropout_Block1[0][0]
Add_FFN_Block1 (Add)	(None, 49, 192)	0	Add_Attn_Block1[0][0], FFN_Projection_Block1[0][0]
Transformer_Norm1_Block2 (LayerNormalization)	(None, 49, 192)	384	Add_FFN_Block1[0][0]
MultiHeadAttention_Block2 (MultiHeadAttention)	(None, 49, 192)	148,224	Transformer_Norm1_Block2[...] Transformer_Norm1_Block2[...]
Attn_Dropout_Block2 (Dropout)	(None, 49, 192)	0	MultiHeadAttention_Block2[...]
Add_Attn_Block2 (Add)	(None, 49, 192)	0	Add_FFN_Block1[0][0], Attn_Dropout_Block2[0][0]
Transformer_Norm2_Block2 (LayerNormalization)	(None, 49, 192)	384	Add_Attn_Block2[0][0]
FFN_Expansion_Block2 (Dense)	(None, 49, 768)	148,224	Transformer_Norm2_Block2[...]
FFN_Dropout_Block2 (Dropout)	(None, 49, 768)	0	FFN_Expansion_Block2[0][0]
FFN_Projection_Block2 (Dense)	(None, 49, 192)	147,648	FFN_Dropout_Block2[0][0]
Add_FFN_Block2 (Add)	(None, 49, 192)	0	Add_Attn_Block2[0][0], FFN_Projection_Block2[0][0]
Transformer_Norm1_Block3 (LayerNormalization)	(None, 49, 192)	384	Add_FFN_Block2[0][0]
MultiHeadAttention_Block3 (MultiHeadAttention)	(None, 49, 192)	148,224	Transformer_Norm1_Block3[...] Transformer_Norm1_Block3[...]
Attn_Dropout_Block3 (Dropout)	(None, 49, 192)	0	MultiHeadAttention_Block3[...]
Add_Attn_Block3 (Add)	(None, 49, 192)	0	Add_FFN_Block2[0][0], Attn_Dropout_Block3[0][0]
Transformer_Norm2_Block3 (LayerNormalization)	(None, 49, 192)	384	Add_Attn_Block3[0][0]
FFN_Expansion_Block3 (Dense)	(None, 49, 768)	148,224	Transformer_Norm2_Block3[...]
FFN_Dropout_Block3 (Dropout)	(None, 49, 768)	0	FFN_Expansion_Block3[0][0]
FFN_Projection_Block3 (Dense)	(None, 49, 192)	147,648	FFN_Dropout_Block3[0][0]
Add_FFN_Block3 (Add)	(None, 49, 192)	0	Add_Attn_Block3[0][0], FFN_Projection_Block3[0][0]
Transformer_Norm1_Block4 (LayerNormalization)	(None, 49, 192)	384	Add_FFN_Block3[0][0]
MultiHeadAttention_Block4 (MultiHeadAttention)	(None, 49, 192)	148,224	Transformer_Norm1_Block4[...] Transformer_Norm1_Block4[...]
Attn_Dropout_Block4 (Dropout)	(None, 49, 192)	0	MultiHeadAttention_Block4[...]
Add_Attn_Block4 (Add)	(None, 49, 192)	0	Add_FFN_Block3[0][0], Attn_Dropout_Block4[0][0]
Transformer_Norm2_Block4 (LayerNormalization)	(None, 49, 192)	384	Add_Attn_Block4[0][0]
FFN_Expansion_Block4 (Dense)	(None, 49, 768)	148,224	Transformer_Norm2_Block4[...]
FFN_Dropout_Block4 (Dropout)	(None, 49, 768)	0	FFN_Expansion_Block4[0][0]
FFN_Projection_Block4 (Dense)	(None, 49, 192)	147,648	FFN_Dropout_Block4[0][0]
Add_FFN_Block4 (Add)	(None, 49, 192)	0	Add_Attn_Block4[0][0], FFN_Projection_Block4[0][0]
Spatial_Sequence_BiGRU (bidirectional)	(None, 256)	247,296	Add_FFN_Block4[0][0]
BiGRU_Dropout (Dropout)	(None, 256)	0	Spatial_Sequence_BiGRU[0][...]
dense_2 (Dense)	(None, 128)	32,896	BiGRU_Dropout[0][0]
batch_normalization (BatchNormalization)	(None, 128)	512	dense_2[0][0]
dropout_4 (Dropout)	(None, 128)	0	batch_normalization[0][0]
dense_3 (Dense)	(None, 64)	8,256	dropout_4[0][0]
batch_normalization_1 (BatchNormalization)	(None, 64)	256	dense_3[0][0]
Output_Probability (Dense)	(None, 1)	65	batch_normalization_1[0][...]

Total params: 2,457,779 (9.38 MB)
Trainable params: 2,457,011 (9.37 MB)
Non-trainable params: 768 (3.00 KB)

To run: `docker run --rm -it --gpus all ``

`-v /usr/lib/wsl:/usr/lib/wsl ``

`-e NVIDIA_VISIBLE_DEVICES=all ``

`-e NVIDIA_DRIVER_CAPABILITIES=all ``

`-e`

`LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" ``

`-e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" ``

`-e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" ``

`-v "${PWD}:/workspace" ``

`tc-dl-pipeline:latest /workspace/scripts/train_vit_gru.py --leadtime 24 --batch-size 36 --workers 8`

9) Train ensemble model using `train_ensemble.py`

network.

Key Architectural Features

- Robust Penultimate Feature Extraction:**
 - Extracts clean, rich 64-dimensional dense embeddings directly from the penultimate layer of both trained models (`best_cnn_lstm_24h.keras` and `best_vit_gru_24h.keras`).
- Trainable Projection & Layer Normalization:**
 - Projects both 64-dim embedding streams to a unified 128-dimensional space followed by `LayerNormalization` to align representation scales.
- Cross-Feature Interaction & Softmax Adaptive Gating:**
 - Computes element-wise multiplication ($\mathbf{v}_{\text{cnn}} \odot \mathbf{v}_{\text{vit}}$) to capture model consensus.
 - Learns sample-specific softmax gate weights ($w_{\text{cnn}}, w_{\text{vit}}$) to dynamically trust whichever backbone is best suited for each individual atmospheric pattern.
- Multi-Head Cross-Attention Fusion:**
 - Shapes the gated embeddings into sequence tokens and processes them through `MultiHeadAttention` to model non-linear interactions across spatial feature representations.
- Deep MLP Head with GELU & Regularization:**
 - Passes the attended representations through a $256 \rightarrow 128 \rightarrow 64$ dense MLP hierarchy with GELU activations, Dropout (0.3), and Batch Normalization.
- PR-AUC Monitoring & AdamW Optimizer:**
 - Trained with `AdamW(learning_rate=1e-4, weight_decay=1e-4)` and monitored via `val_pr_auc` for early stopping.

🔑 Key Engineering & Architectural Fixes Included

1. **Robust Layer Name Harvesting (`get_layer`):**
 - Eliminates positional indexing entirely. Explicitly queries target layer names from both backbones (`Row_LSTM` , `Col_LSTM` , `Spatial_Sequence_BiGRU` , and `Dense` layers).
2. **Rich 320-Dimensional Backbone Representations:**
 - Concatenates recurrent sequence states (LSTM/BiGRU) with both intermediate dense layers (128 + 64 + 128) to preserve 320 total structural features per model (640 total raw backbone dimensions) instead of collapsing them prematurely.
3. **Probability & Calibration Feature Branch:**
 - Explicitly extracts scalar outputs (p_{cnn} , p_{vit}) and computes a 5-dimensional probability calibration vector:

$$\mathbf{v}_{\text{prob}} = \left[p_{\text{cnn}}, p_{\text{vit}}, |p_{\text{cnn}} - p_{\text{vit}}|, \frac{p_{\text{cnn}} + p_{\text{vit}}}{2}, p_{\text{cnn}} \cdot p_{\text{vit}} \right]$$

4. **16-Chunk Sequence Projection for Multi-Head Attention:**
 - Projects and splits the 320-dim feature spaces of both models into 8 spatial chunk tokens each, creating a rich **16-token sequence** (16×64) for true non-linear cross-attention (16×16 attention matrix).
5. **Learnable Residual Scaling (α) & DropPath:**
 - Implements LayerScale (α initialized to 0.1) and LayerNorm over interaction vectors ($\mathbf{v}_{\text{prod}}$, $\mathbf{v}_{\text{diff}}$) to prevent gradient dominance and overfitting.
6. **Scalar Confidence Gating & Compact MLP Head:**
 - Learns scalar confidence multipliers (g_{cnn} , $g_{\text{vit}} \in [0, 1]$ via Sigmoid) and processes fused features through a tight $128 \rightarrow 64 \rightarrow 1$ MLP head with `BinaryCrossentropy(label_smoothing=0.02)`.
7. **Two-Stage Fine-Tuning Execution:**
 - **Stage 1 (Epochs 1–15):** Trains the fusion head with frozen backbones ($\text{lr} = 5 \times 10^{-5}$).
 - **Stage 2 (Epochs 16–35):** Unfreezes the top layers of both backbones for joint end-to-end fine-tuning at $\text{lr} = 10^{-6}$.

2026-07-24 07:02:01.510547: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1928] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 2267 MB memory
Model: "Advanced_16Token_Gated_Ensemble"

Layer (type)	Output Shape	Param #	Connected to
Unified_Input_Grid (InputLayer)	(None, 21, 21, 137)	0	-
CNN_Harvester (Functional)	[(None, 320), (None, 1)]	540,161	Unified_Input_Grid[0][0]
ViT_Harvester (Functional)	[(None, 320), (None, 1)]	2,490,675	Unified_Input_Grid[0][0]
Proj_CNN_512 (Dense)	(None, 512)	164,352	CNN_Harvester[0][0]
Proj_ViT_512 (Dense)	(None, 512)	164,352	ViT_Harvester[0][0]
Norm_CNN_512 (LayerNormalization)	(None, 512)	1,024	Proj_CNN_512[0][0]
Norm_ViT_512 (LayerNormalization)	(None, 512)	1,024	Proj_ViT_512[0][0]
Prob_Diff (Lambda)	(None, 1)	0	CNN_Harvester[0][1], ViT_Harvester[0][1]
Prob_Mean (Lambda)	(None, 1)	0	CNN_Harvester[0][1], ViT_Harvester[0][1]
Prob_Prod (Multiply)	(None, 1)	0	CNN_Harvester[0][1], ViT_Harvester[0][1]
reshape_1 (Reshape)	(None, 8, 64)	0	Norm_CNN_512[0][0]
reshape_2 (Reshape)	(None, 8, 64)	0	Norm_ViT_512[0][0]
Probability_Calibration_Vect_ (Concatenate)	(None, 5)	0	CNN_Harvester[0][1], ViT_Harvester[0][1], Prob_Diff[0][0], Prob_Mean[0][0], Prob_Prod[0][0]
global_average_pooling1d (GlobalAveragePooling1D)	(None, 64)	0	reshape_1[0][0]
global_average_pooling1d_1 (GlobalAveragePooling1D)	(None, 64)	0	reshape_2[0][0]
Gate_Context (Concatenate)	(None, 133)	0	Probability_Calibration_Vect_, global_average_pooling1d_1, global_average_pooling1d_1
Scalar_Gate_CNN (Dense)	(None, 1)	134	Gate_Context[0][0]
Scalar_Gate_ViT (Dense)	(None, 1)	134	Gate_Context[0][0]
Gated_CNN_512 (Multiply)	(None, 512)	0	Norm_CNN_512[0][0], Scalar_Gate_CNN[0][0]
Gated_ViT_512 (Multiply)	(None, 512)	0	Norm_ViT_512[0][0], Scalar_Gate_ViT[0][0]
Tokens_CNN (Reshape)	(None, 8, 64)	0	Gated_CNN_512[0][0]
Tokens_ViT (Reshape)	(None, 8, 64)	0	Gated_ViT_512[0][0]
Tokens_16_Sequence (Concatenate)	(None, 16, 64)	0	Tokens_CNN[0][0], Tokens_ViT[0][0]
Cross_Attention_16_Tokens (MultiHeadAttention)	(None, 16, 64)	16,640	Tokens_16_Sequence[0][0], Tokens_16_Sequence[0][0]
LayerScale_Attn (LayerScale)	(None, 16, 64)	64	Cross_Attention_16_Tokens_16_Tokens
Interaction_Subtract (Subtract)	(None, 512)	0	Norm_CNN_512[0][0], Norm_ViT_512[0][0]
add (Add)	(None, 16, 64)	0	Tokens_16_Sequence[0][0], LayerScale_Attn[0][0]
Interaction_Product (Multiply)	(None, 512)	0	Norm_CNN_512[0][0], Norm_ViT_512[0][0]
Interaction_AbsDiff (Lambda)	(None, 512)	0	Interaction_Subtract[0][0]
Norm_Attn_Residual (LayerNormalization)	(None, 16, 64)	128	add[0][0]
Norm_Prod (LayerNormalization)	(None, 512)	1,024	Interaction_Product[0][0]
Norm_Diff (LayerNormalization)	(None, 512)	1,024	Interaction_AbsDiff[0][0]
Global_Token_Pool (GlobalAveragePooling1D)	(None, 64)	0	Norm_Attn_Residual[0][0]
Prod_Compress (Dense)	(None, 64)	32,832	Norm_Prod[0][0]
Diff_Compress (Dense)	(None, 64)	32,832	Norm_Diff[0][0]
Master_Fusion_Vector (Concatenate)	(None, 197)	0	Global_Token_Pool[0][0], Prod_Compress[0][0], Diff_Compress[0][0], Probability_Calibration_Vect_
MLP_Dense128 (Dense)	(None, 128)	25,344	Master_Fusion_Vector[0][0]
MLP_BN1 (BatchNormalization)	(None, 128)	512	MLP_Dense128[0][0]
MLP_GELU1 (Activation)	(None, 128)	0	MLP_BN1[0][0]
MLP_Dropout1 (Dropout)	(None, 128)	0	MLP_GELU1[0][0]
MLP_Dense64 (Dense)	(None, 64)	8,256	MLP_Dropout1[0][0]
MLP_BN2 (BatchNormalization)	(None, 64)	256	MLP_Dense64[0][0]
MLP_GELU2 (Activation)	(None, 64)	0	MLP_BN2[0][0]
MLP_Dropout2 (Dropout)	(None, 64)	0	MLP_GELU2[0][0]
Ensemble_Probability (Dense)	(None, 1)	65	MLP_Dropout2[0][0]

Total params: 3,480,833 (13.28 MB)
Trainable params: 482,509 (1.84 MB)
Non-trainable params: 2,998,324 (11.44 MB)

10)evaluating model performance on test set

Let's run the evaluation on your holdout test set (`tc_24h_test.csv`) to see how the model generalizes to completely unseen storm seasons (2021–2022).

`scripts/evaluate_model.py`:for cnn+vit

`scripts/evaluate_ensemble.py` for ensemble

11)script to generate full test classification log file (`test_predictions_log.csv`) so we can review the actual vs. predicted labels.

`log_test_predictions.py`:

Key Upgrades Included:

1. **Multi-Model Predictions:** Logs actual vs. predicted probabilities for CNN-LSTM, ViT-BiGRU, and the Gated Cross-Attention Ensemble side-by-side in a single CSV file.
2. **Safe Weight-Loading Architecture:** Uses graph reconstruction via `build_advanced_ensemble()` and direct weight loading (`.load_weights()`) to completely bypass Keras 3 custom layer deserialization errors.
3. **Flexible CLI Arguments:** Accepts `--leadtime` (24, 48, 72h) to automatically load the correct dataset splits and checkpoints.
4. **Vectorized Inference Pass:** Performs batch prediction over the entire dataset generator at once instead of calling `.predict()` in a slow row-by-row loop.


Classification log generated successfully at:

`data/extracted_features_entire_world/test_predictions_log_24h.csv`

12)Geographical Visualization:

Plot Target_Latitude and Target_Longitude on a global map using cartopy or plotly to visualize model performance across ocean basins on test set(Atlantic, East Pacific, West Pacific, Indian Ocean).

`scripts/plot_spatial_performance.py`



13)XAI

SHOW CODE <>

Yes! The updated script generates all 5 XAI outputs.

`scripts/explain_ensemble_24h.py` generates 5 distinct, publication-ready figures saved directly inside `plots/xai/`:

The 5 Generated Outputs


1. `1_gradcam_heatmap_24h.png` (**Grad-CAM**):
Displays a 2D spatial heatmap overlay (21×21 patch) showing where the CNN feature extractor detects spatial convective organization.
2. `2_integrated_gradients_map_24h.png` (**Integrated Gradients**):
Computes exact gradient-based saliency across all 137 atmospheric channels to highlight the storm's core circulation center.
3. `3_vit_attention_rollout_24h.png` (**ViT Attention Rollout**):
Renders the self-attention token weights from the Vision Transformer encoder patches to map spatial dependencies across the 21×21 region.
4. `4_grib_variable_importance_24h.png` (**GRIB Variable Importance**):
Ranks your **11 environmental variables** (`absvprs` , `capescfc` , `rhprs` , `vvelprs` , `ugrdprs` , `vgrdprs` , `tmpprs` , `hgtpsr` , `pressfc` , `tmprsc` , `landmask`) by their relative percentage contribution to triggering a genesis prediction.
5. `5_vertical_pressure_profile_24h.png` (**Vertical Pressure Profile**):
Plots attribution scores across all **19 vertical pressure levels** (1000 hPa down to 200 hPa) to show which atmospheric altitude (e.g., 850 hPa low-level vorticity vs 200 hPa upper divergence) drives storm formation.

 Save & Replace `scripts/explain_ensemble_24h.py`

Stage 5 XAI Roadmap for Your Keras/TF Pipeline

We will generate 5 key interpretability outputs:

1. **Grad-CAM (Spatial Heatmaps):** Applied to the final convolutional/feature layer of your trained CNN backbone to highlight spatial regions (21×21 patches) driving genesis predictions.
2. **Integrated Gradients (Pixel & Channel Saliency):** Computes exact gradient-based attributions across both spatial dimensions and feature channels.
3. **ViT Attention Rollout:** Extracts multi-head self-attention weights from your Vision Transformer patches to map global token interactions.
4. **SHAP / Permutation Feature Importance:** Ranks your **11 extracted GRIB environmental variables** to show which meteorological parameters (e.g., `absvprs` , `capescfc` , `rhprs` , `vvelprs`) trigger the genesis decision.
5. **Vertical Profile Attribution:** Collapses pressure levels (1000 hPa – 200 hPa) to show which atmospheric height (e.g., 850 hPa low-level spin vs 200 hPa upper divergence) is most important.





✅ Feature importance chart saved to: plots/xai/grib_feature_importance_24h.png


===== FEATURE IMPORTANCE RANKING =====

```
1. absvprs      : 19.86%
2. rhprs        : 17.03%
3. vvelprs      : 15.26%
4. ugrdprs      : 15.07%
5. vgrdprs      : 11.01%
6. hgtpers      : 10.21%
7. tmpprs       : 9.22%
8. capesfc      : 0.85%
9. tmpsfc       : 0.72%
10. pressfc     : 0.64%
11. landmask    : 0.13%
```


PS E:\TC-DL>



2. Integrated Gradients Saliency Map (24h Lead Time)



3. ViT Patch Attention Rollout Map (24h Lead Time)



Here is a structured analysis of the explainability and feature attribution results for your **24-hour lead time** tropical cyclone prediction model across spatial, atmospheric, and vertical dimensions.

1. Spatial Explainability & Feature Saliency

Interpretability Method	Primary Focus Region (<i>x, y</i>)	Spatial Characteristics & Behavior
Grad-CAM	Coarse peak at (16, 10)	Highlights region-level feature maps in upper CNN/backbone layers; strong gradient energy toward the eastern/right quadrant.
Integrated Gradients	Tightly focused at (13, 9)	High-resolution, pixel-level attribution mapping; pinpoints the exact dynamic center/vorticity core driving the forecast.
ViT Attention Rollout	Broad swath across (12–20, 7–13)	Long-range context aggregation; captures spatial interactions and atmospheric environment surrounding the storm center.

...

- **Synthesis:** The spatial alignment across local gradients (IG), CNN feature maps (Grad-CAM), and global self-attention (ViT) confirms that the model is actively tracking the cyclone core and its surrounding circulation environment rather than relying on peripheral noise or boundary artifacts.

2. Environmental Variable Importance

The relative feature attribution scores strongly favor **pressure-level (multilevel) variables** over surface-level metrics:

- **Primary Drivers (> 15% score):**
 - **Absolute Vorticity** (`absvprs` , ~ 19.6%): The single most dominant factor, driving storm circulation dynamics.
 - **Relative Humidity** (`rhprs` , ~ 17.1%): Crucial for moisture flux and eyewall convection maintenance.
 - **Vertical Velocity** (`vvelprs` , ~ 15.6%): Measures updraft/downdrafts and convective instability.
 - **U-Component Wind** (`ugrdprs` , ~ 15.3%): Zonal wind vector driving steering flow and horizontal shear.
- **Secondary Atmospheric Drivers (9%–11%):**
 - **V-Component Wind** (`vgrdprs` , ~ 11.0%)
 - **Geopotential Height** (`hgtprs` , ~ 10.1%)
 - **Temperature Profile** (`tmpprs` , ~ 9.0%)
- **Surface & Land Mask Contributions (< 1% each):**
 - **Surface CAPE** (`capesfc`), **Surface Temp** (`tmptsfc`), **Surface Pressure** (`pressfc`), and `landmask` contribute minimally at 24h lead time relative to the atmospheric column.

- o Surface CAPE (`capesfc`), Surface Temp (`tmpsfc`), Surface Pressure (`pressfc`), and `landmask` contribute minimally at 24h lead time relative to the atmospheric column.

3. Vertical Pressure Profile Structure

The vertical attribution distribution spans pressure levels from 1000 hPa (surface) to 200 hPa (tropopause):

- o **Peak Attribution (450 hPa):** Highest importance reaches $\sim 6.22\%$ at mid-troposphere (450 hPa), where deep convection and latent heat release strongly govern mid-level vortex development.
- o **Lower-to-Mid Tropospheric Band (800–550 hPa):** Consistently elevated contribution ($\sim 5.5\%$ – 5.9%), reflecting low-to-mid level inflow and momentum transport.
- o **Boundary Layer & Upper Troposphere:** Lower attribution near the surface (1000–950 hPa at $\sim 4.5\%$ – 4.6%) and upper troposphere (250–200 hPa at $\sim 5.0\%$ – 5.5%).

Summary Findings

The 24h lead-time model has learned meteorologically consistent representations: it grounds its predictions primarily in **mid-tropospheric absolute vorticity, moisture availability, and vertical motion (800–450 hPa)** centered tightly around the cyclone's dynamical core.

👍 🗨️ ↻ 📄 ...

14)generate_runtime_configs.py

Updated the script so it can **generate the validation prediction CSV itself**, then calculate and save the threshold files. [Share](#) ...

🔗 `generate_runtime_configs.py`

It will generate:

```
data/extracted_features_entire_world/validation_predictions_log_24h.csv
saved_models/training_metadata.json
saved_models/operational_thresholds.json
saved_models/threshold_selection_report.json
```

Run from `E:\TC-DL`:

```
PowerShell
docker run --rm -it --gpus all `
-v /usr/lib/wsl:/usr/lib/wsl `
-e NVIDIA_VISIBLE_DEVICES=all `
-e NVIDIA_DRIVER_CAPABILITIES=all `
-e LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn" `
-e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
-e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
-v "${PWD}:/workspace" `
tc-dl-pipeline:latest `
/workspace/generate_runtime_configs.py `
--generate-validation `
--leadtime 24 `
--validation-csv /workspace/data/extracted_features_entire_world/tc_24h_val.csv `
--features-dir /workspace/data/extracted_features_entire_world `
--batch-size 32 `
--selection_min_recall `
--target-recall 0.90
```

The generated CSV will use the actual validation labels and fresh ensemble probabilities, so the operational threshold is selected without touching the holdout test set.

Your existing test log uses `Actual_Label` and `Prob_Ensemble`, confirming the same output format expected for the validation log. [Pasted text](#) `</>`

```

docker run --rm -it --gpus all `
-v /usr/lib/wsl:/usr/lib/wsl `
-e NVIDIA_VISIBLE_DEVICES=all `
-e NVIDIA_DRIVER_CAPABILITIES=all `
-e
LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" `
-e http_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
-e https_proxy="http://rrscnorth:NRSC%40User@192.168.0.9:8080" `
-v "${PWD}:/workspace" `
tc-dl-pipeline:latest `
/workspace/scripts/generate_validation_and_configs.py `
--leadtime 24 `
--validation-csv /workspace/data/extracted_features_entire_world/tc_24h_val.csv `
--features-dir /workspace/data/extracted_features_entire_world `
--batch-size 32 `
--selection min_recall `
--target-recall 0.90

```

14)run_fnl_2026_cyclogenesis_verification.py

```

docker run --rm -it --gpus all -v /usr/lib/wsl:/usr/lib/wsl -e NVIDIA_VISIBLE_DEVICES=all -e
NVIDIA_DRIVER_CAPABILITIES=all -e
LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" -v "${PWD}:/workspace" tc-dl-pipeline:latest
/workspace/scripts/run_fnl_2026_cyclogenesis_verification.py --leadtime 24 --ibtracs-csv
/workspace/data/ibtracs/ibtracs.ALL.list.v04r01.csv --start-date 2026-01-01 --end-date 2027-01-01
--genesis-definition first-track --lead-tolerance-hours 0 --alert-threshold 0.75 --min-expert-confidence
0.45 --minimum-midlevel-rh-percent 50 --minimum-vorticity-850-s1 7.5e-6
--maximum-pressure-anomaly-pa 0 --maximum-shear-200-850-ms 25
--maximum-lowlevel-divergence-s1=1.0e-4 --min-cluster-support 2 --cluster-radius-km 700
--max-detections 2 --ibtracs-distance-tolerance-km 800 --batch-size 256 --download-workers 4
--insecure-ssl

```

15)main.py

```

docker run --rm -it --gpus all `

```

```
-v /usr/lib/wsl:/usr/lib/wsl `
-e NVIDIA_VISIBLE_DEVICES=all `
-e NVIDIA_DRIVER_CAPABILITIES=all `
-e
LD_LIBRARY_PATH="/usr/lib/wsl/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib"
`
-v "${PWD}:/workspace" `
tc-dl-pipeline:latest `
/workspace/main.py `
--once `
--lead-times 24 `
--alert-threshold 0.75 `
--min-expert-confidence 0.45 `
--nms-radius-km 500 `
--max-detections 5 `
--batch-size 256 `
--classification-mode auto `
--post-classifier-model
/workspace/saved_models/genesis_post_classifier.json `
--active-systems-json /workspace/data/active_systems.json `
--insecure-ssl
```

16)update_active_sytems.py

```
docker run --rm -it `
-v "${PWD}:/workspace" `
--entrypoint python3 `
```

```
tc-dl-pipeline:latest `
/workspace/scripts/update_active_systems.py `
--output /workspace/data/active_systems.json `
--max-age-hours 72 `
--allow-partial `
--insecure-ssl
```

17)index.html

```
docker run --rm -it `
-p 8080:8080 `
-v "${PWD}\output:/site:ro" `
--entrypoint python3 `
tc-dl-pipeline:latest `
-m http.server 8080 `
--directory /site
```

18)XAI POST CLASSIFIER

```
build_post_classifier_fast.py
E:\TC-DL\scripts\build_post_classifier_fast.py
```

Purpose:

- finds existing historical prediction CSV files,
- runs additional historical verification,
- collects HIT and FALSE_ALARM rows,
- stops after enough samples are collected,
- calculates environmental features,

- calculates batched XAI,
- trains logistic regression,
- compares it against fallback rules,
- writes the learned JSON only if it performs better.

```
docker run --rm -it --gpus all `
```

```
>> -v /usr/lib/wsl:/usr/lib/wsl `
```

```
>> -e NVIDIA_VISIBLE_DEVICES=all `
```

```
>> -e NVIDIA_DRIVER_CAPABILITIES=all `
```

```
>> -e
```

```
LD_LIBRARY_PATH="/usr/lib/wsl/lib/./usr/local/lib/python3.10/dist-packages/nvidia/cudnn/lib:/usr/local/lib/python3.10/dist-packages/nvidia/cublas/lib" `
```

```
>> -v "${PWD}:/workspace" `
```

```
>> --entrypoint python3 `
```

```
>> tc-dl-pipeline:latest `
```

```
>> /workspace/scripts/build_post_classifier_fast.py `
```

```
>> --verifier-script /workspace/scripts/run_fnl_2026_cyclogenesis_verification.py `
```

```
>> --target-samples 100 `
```

```
>> --minimum-hits 20 `
```

```
>> --minimum-false-alarms 40 `
```

```
>> --start-year 2025 `
```

```
>> --start-month 7 `
```

```
>> --months 5 `
```

```
>> --batch-size 256 `
```

```
>> --xai-batch-size 16 `
```

```
>> --download-workers 4 `
```

```
>> --insecure-ssl
```

E:\TC-DL\saved_models\genesis_post_classifier.json

Purpose:

- stores the fitted post-classifier,
- does not require TensorFlow or scikit-learn during operational classification,
- allows `main.py` to calculate the logistic probability directly.

It contains fields similar to:

```
{  
  "model_type": "logistic_regression",  
  "uses_xai": true,  
  "feature_names": [],  
  "coefficients": {},  
  "intercept": 0.0,  
  "means": {},  
  "scales": {},  
  "imputation_medians": {},  
  "positive_threshold": 0.25,
```

```
"high_priority_threshold": 0.40  
}
```

Combined training CSV

**E:\TC-DL\data\post_classifier_training\
genesis_post_classifier_training.csv**

Purpose:

- stores all post-classifier training rows,
- contains labels,
- CNN/ViT/ensemble scores,
- meteorological diagnostics,
- XAI values,
- matched cyclone information.

This is the main dataset used to train the JSON classifier.